

detectron2-onnx

April 9, 2021

1 Detectron2 onnx export and inference

To gain some speed, especially in embedded devices such as the raspberry pi. Deploying a detectron2 model and run the inference on a CPU based backend like the caffe2 onnx backend is a must.

This notebooks describes how to do this. It is based on the [official colab](#).

Run this notebook inside the detectron2 repository

2 Install detectron2

```
[1]: import torch, torchvision
print(torch.__version__, torch.cuda.is_available())
```

1.8.0 False

```
[2]: # Some basic setup:
# Setup detectron2 logger
import detectron2
from detectron2.utils.logger import setup_logger
setup_logger()

import time

# import some common libraries
import numpy as np
import os, json, cv2, random
import matplotlib.pyplot as plt

plt.rcParams['figure.dpi'] = 200

def cv2_imshow(im):
    plt.imshow(cv2.cvtColor(im, cv2.COLOR_BGR2RGB))

# import some common detectron2 utilities
from detectron2 import model_zoo
from detectron2.engine import DefaultPredictor
from detectron2.config import get_cfg
```

```

from detectron2.utils.visualizer import Visualizer
from detectron2.data import MetadataCatalog, DatasetCatalog

from detectron2.data import build_detection_test_loader

import detectron2.data.transforms as T

```

3 Run a pre-trained detectron2 model

We first download an image from the COCO dataset:

```

[ ]: # prepare the coco test dataset
!wget http://images.cocodataset.org/zips/val2017.zip
!mkdir -p datasets/coco/
!unzip val2017.zip -d datasets/coco/
./datasets/prepare_for_tests.sh
!mv datasets/coco/annotations/instances_val2017_100.json datasets/coco/
↪ annotations/instances_val2017.json

```

Create a detectron2 config. Choose any model from the [model zoo](#).

```

[3]: model_config = "COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml"

[4]: cfg = get_cfg()

# add project-specific config (e.g., TensorMask) here if you're not running a
↪ model in detectron2's core library
cfg.merge_from_file(model_zoo.get_config_file(model_config))
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5 # set threshold for this model

# Find a model from detectron2's model zoo. You can use the https://dl.
↪ fbaipublicfiles... url as well
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url(model_config)
cfg.MODEL.DEVICE = 'cpu'

[5]: # Load an image from the coco dataset.
original_image = cv2.imread('datasets/coco/val2017/000000023937.jpg')

[6]: # Use the detectron2 `DefaultPredictor` to run inference on this image.
predictor = DefaultPredictor(cfg)
outputs = predictor(original_image)

[7]: # We can use `Visualizer` to draw the predictions on the image.
v = Visualizer(original_image[:, :, ::-1], MetadataCatalog.get(cfg.DATASETS.
↪ TRAIN[0]), scale=1.2)
out = v.draw_instance_predictions(outputs["instances"].to("cpu"))

```

```
cv2_imshow(out.get_image()[:, :, ::-1])
```



4 Export the model as onnx

Detectron2 already provides a script to export onnx models, use it with:

```
./tools/deploy/export_model.py --config-file configs/COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x/137849600/model.py \
    --output ./output \
    --export-method caffe2_tracing \
    --format onnx \
    MODEL.WEIGHTS detectron2://COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x/137849600/model.py \
    MODEL.DEVICE cpu
```

The following steps show how the onnx tracing is done.

```
[8]: from detectron2.checkpoint import DetectionCheckpointer
      from detectron2.modeling import build_model

      from detectron2.export import Caffe2Tracer

      from torchvision.transforms import transforms

      import onnx
```

```
[9]: # build the model with the config used before
      torch_model = build_model(cfg)
```

```
DetectionCheckpointer(torch_model).resume_or_load(cfg.MODEL.WEIGHTS)
torch_model.eval()
```

```
[9]: GeneralizedRCNN(
  (backbone): FPN(
    (fpn_lateral2): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (fpn_lateral3): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (fpn_lateral4): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (fpn_lateral5): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output5): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (top_block): LastLevelMaxPool()
    (bottom_up): ResNet(
      (stem): BasicStem(
        (conv1): Conv2d(
          3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False
        )
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
    )
    (res2): Sequential(
      (0): BottleneckBlock(
        (shortcut): Conv2d(
          64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        )
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv1): Conv2d(
        64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
    )
    (conv2): Conv2d(
      64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
    )
    (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
  )
  (conv3): Conv2d(
    64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
  )
  (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
)
  (1): BottleneckBlock(
    (conv1): Conv2d(
```

```

        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
    )
    (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
    )
    (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
)
(2): BottleneckBlock(
    (conv1): Conv2d(
        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
    )
    (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
    )
    (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
)
)
(res3): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
        256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv1): Conv2d(
        256, 128, kernel_size=(1, 1), stride=(2, 2), bias=False
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv2): Conv2d(
        128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    )
    (conv3): Conv2d(
        128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)

```

```

    )
)
(1): BottleneckBlock(
  (conv1): Conv2d(
    512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv2): Conv2d(
    128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv3): Conv2d(
    128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
  )
)
(2): BottleneckBlock(
  (conv1): Conv2d(
    512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv2): Conv2d(
    128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv3): Conv2d(
    128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
  )
)
(3): BottleneckBlock(
  (conv1): Conv2d(
    512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv2): Conv2d(
    128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv3): Conv2d(
    128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
  )
)
)

```

```

)
(res4): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
      512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
    (conv1): Conv2d(
      512, 256, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (1): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (2): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(

```

```

        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
(3): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(4): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(5): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)

```



```

    )
    )
    )
    (res5): Sequential(
      (0): BottleneckBlock(
        (shortcut): Conv2d(
          1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False
          (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
        )
        (conv1): Conv2d(
          1024, 512, kernel_size=(1, 1), stride=(2, 2), bias=False
          (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
        )
        (conv2): Conv2d(
          512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
          (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
        )
        (conv3): Conv2d(
          512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
          (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
        )
      )
      (1): BottleneckBlock(
        (conv1): Conv2d(
          2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
          (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
        )
        (conv2): Conv2d(
          512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
          (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
        )
        (conv3): Conv2d(
          512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
          (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
        )
      )
      (2): BottleneckBlock(
        (conv1): Conv2d(
          2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
          (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
        )
        (conv2): Conv2d(
          512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
          (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)

```



```
[10]: # https://github.com/facebookresearch/detectron2/blob/master/tools/deploy/
      ↪ export_model.py#L45
      from detectron2.data import build_detection_test_loader

      # get a sample data
      data_loader = build_detection_test_loader(cfg, cfg.DATASETS.TEST[0])
      first_batch = next(iter(data_loader))

      # https://github.com/facebookresearch/detectron2/blob/master/tools/deploy/
      ↪ export_model.py#L45
      tracer = Caffe2Tracer(
          cfg,
          torch_model,
          first_batch
      )

      onnx_model = tracer.export_onnx()
      onnx.save(onnx_model, "model.onnx")
```

```
[03/18 19:26:47 d2.data.datasets.coco]: Loaded 100 images in COCO
format from datasets/coco/annotations/instances_val2017.json
[03/18 19:26:47 d2.data.build]: Distribution of instances among all 80
categories:
```

category	#instances	category	#instances	category	#instances
person	341	bicycle	10	car	51
motorcycle	23	airplane	0	bus	10
train	2	truck	4	boat	13
traffic light	9	fire hydrant	5	stop sign	1
parking meter	0	bench	14	bird	13
cat	2	dog	4	horse	5
sheep	0	cow	1	elephant	3
bear	6	zebra	16	giraffe	2
backpack	8	umbrella	8	handbag	12
tie	11	suitcase	0	frisbee	9
skis	8	snowboard	1	sports ball	4
kite	19	baseball bat	2	baseball glove	1
skateboard	1	surfboard	3	tennis racket	9
bottle	16	wine glass	5	cup	15
fork	2	knife	0	spoon	1
bowl	7	banana	1	apple	2

```

[03/18 19:26:47 d2.data.dataset_mapper]: [DatasetMapper] Augmentations
used in inference: [ResizeShortestEdge(short_edge_length=(800, 800),
max_size=1333, sample_style='choice')]
[03/18 19:26:47 d2.data.common]: Serializing 100 elements to byte
tensors and concatenating them all ...
[03/18 19:26:47 d2.data.common]: Serialized dataset takes 0.46 MiB

/home/users/fassbinb/.local/lib/python3.8/site-packages/torch/onnx/utils.py:262:
UserWarning: `add_node_names` can be set to True only when
`operator_export_type` is `ONNX`. Since `operator_export_type` is not set to
`ONNX`, `add_node_names` argument will be ignored.
  warnings.warn("`{}' can be set to True only when `operator_export_type` is ")
/home/users/fassbinb/.local/lib/python3.8/site-packages/torch/onnx/utils.py:262:
UserWarning: `do_constant_folding` can be set to True only when
`operator_export_type` is `ONNX`. Since `operator_export_type` is not set to
`ONNX`, `do_constant_folding` argument will be ignored.
  warnings.warn("`{}' can be set to True only when `operator_export_type` is ")
/home/users/fassbinb/Work/detectron2/detectron2/export/c10.py:31: TracerWarning:
Converting a tensor to a Python boolean might cause the trace to be incorrect.
We can't record the data flow of Python values, so this value will be treated as
a constant in the future. This means that the trace might not generalize to
other inputs!
  assert tensor.dim() == 2 and tensor.size(-1) in [4, 5, 6], tensor.size()
/home/users/fassbinb/Work/detectron2/detectron2/export/c10.py:385:
TracerWarning: Converting a tensor to a Python boolean might cause the trace to
be incorrect. We can't record the data flow of Python values, so this value will
be treated as a constant in the future. This means that the trace might not
generalize to other inputs!
  if num_classes + 1 == class_logits.shape[1]:
/home/users/fassbinb/Work/detectron2/detectron2/export/c10.py:394:
TracerWarning: Converting a tensor to a Python boolean might cause the trace to
be incorrect. We can't record the data flow of Python values, so this value will
be treated as a constant in the future. This means that the trace might not
generalize to other inputs!
  assert box_regression.shape[1] % box_dim == 0
/home/users/fassbinb/Work/detectron2/detectron2/export/c10.py:401:
TracerWarning: Converting a tensor to a Python boolean might cause the trace to
be incorrect. We can't record the data flow of Python values, so this value will
be treated as a constant in the future. This means that the trace might not
generalize to other inputs!
  if input_tensor_mode:
/home/users/fassbinb/Work/detectron2/detectron2/export/c10.py:433:
TracerWarning: Converting a tensor to a Python boolean might cause the trace to
be incorrect. We can't record the data flow of Python values, so this value will
be treated as a constant in the future. This means that the trace might not
generalize to other inputs!
  nms_outputs = torch.ops._caffe2.BoxWithNMSLimit(
/home/users/fassbinb/.local/lib/python3.8/site-packages/torch/tensor.py:587:

```

RuntimeWarning: Iterating over a tensor might cause the trace to be incorrect.
Passing a tensor of different shape won't change the number of iterations
executed (and might lead to errors or silently give incorrect results).

```
warnings.warn('Iterating over a tensor might cause the trace to be incorrect.',
```

```
/home/users/fassbinb/Work/detectron2/detectron2/export/c10.py:462:
```

TracerWarning: Converting a tensor to a Python number might cause the trace to be incorrect. We can't record the data flow of Python values, so this value will be treated as a constant in the future. This means that the trace might not generalize to other inputs!

```
for i, b in enumerate(int(x.item()) for x in roi_batch_splits_nms)
```

```
[11]: # the model can be loaded using the Caffe2 backend (optimized for cpu
      ↪ inference) and
      # do predictions without detectron2
      import caffe2.python.onnx.backend as backend

      model = onnx.load("model.onnx") # or use the `onnx_model` from above
      onnx.checker.check_model(model)

      # To print a human readable representation of the graph use:
      print(onnx.helper.printable_graph(model.graph))
```

```
graph torch-jit-export (
  %data[UINT8, 1x3x800x1216]
  %im_info[FLOAT, 1x3]
) optional inputs with matching initializers (
  %_wrapped_model.backbone.fpn_lateral2.weight[FLOAT, 256x256x1x1]
  %_wrapped_model.backbone.fpn_lateral2.bias[FLOAT, 256]
  %_wrapped_model.backbone.fpn_output2.weight[FLOAT, 256x256x3x3]
  %_wrapped_model.backbone.fpn_output2.bias[FLOAT, 256]
  %_wrapped_model.backbone.fpn_lateral3.weight[FLOAT, 256x512x1x1]
  %_wrapped_model.backbone.fpn_lateral3.bias[FLOAT, 256]
  %_wrapped_model.backbone.fpn_output3.weight[FLOAT, 256x256x3x3]
  %_wrapped_model.backbone.fpn_output3.bias[FLOAT, 256]
  %_wrapped_model.backbone.fpn_lateral4.weight[FLOAT, 256x1024x1x1]
  %_wrapped_model.backbone.fpn_lateral4.bias[FLOAT, 256]
  %_wrapped_model.backbone.fpn_output4.weight[FLOAT, 256x256x3x3]
  %_wrapped_model.backbone.fpn_output4.bias[FLOAT, 256]
  %_wrapped_model.backbone.fpn_lateral5.weight[FLOAT, 256x2048x1x1]
  %_wrapped_model.backbone.fpn_lateral5.bias[FLOAT, 256]
  %_wrapped_model.backbone.fpn_output5.weight[FLOAT, 256x256x3x3]
  %_wrapped_model.backbone.fpn_output5.bias[FLOAT, 256]
  %_wrapped_model.proposal_generator.rpn_head.conv.weight[FLOAT, 256x256x3x3]
  %_wrapped_model.proposal_generator.rpn_head.conv.bias[FLOAT, 256]
  %_wrapped_model.proposal_generator.rpn_head.objectness_logits.weight[FLOAT,
3x256x1x1]
  %_wrapped_model.proposal_generator.rpn_head.objectness_logits.bias[FLOAT, 3]
```

```

    %_wrapped_model.proposal_generator.rpn_head.anchor_deltas.weight[FLOAT,
12x256x1x1]
    %_wrapped_model.proposal_generator.rpn_head.anchor_deltas.bias[FLOAT, 12]
    %_wrapped_model.proposal_generator.anchor_generator.cell_anchors.0[FLOAT, 3x4]
    %_wrapped_model.proposal_generator.anchor_generator.cell_anchors.1[FLOAT, 3x4]
    %_wrapped_model.proposal_generator.anchor_generator.cell_anchors.2[FLOAT, 3x4]
    %_wrapped_model.proposal_generator.anchor_generator.cell_anchors.3[FLOAT, 3x4]
    %_wrapped_model.proposal_generator.anchor_generator.cell_anchors.4[FLOAT, 3x4]
    %_wrapped_model.roi_heads.box_head.fc1.weight[FLOAT, 1024x12544]
    %_wrapped_model.roi_heads.box_head.fc1.bias[FLOAT, 1024]
    %_wrapped_model.roi_heads.box_head.fc2.weight[FLOAT, 1024x1024]
    %_wrapped_model.roi_heads.box_head.fc2.bias[FLOAT, 1024]
    %_wrapped_model.roi_heads.box_predictor.cls_score.weight[FLOAT, 81x1024]
    %_wrapped_model.roi_heads.box_predictor.cls_score.bias[FLOAT, 81]
    %_wrapped_model.roi_heads.box_predictor.bbox_pred.weight[FLOAT, 320x1024]
    %_wrapped_model.roi_heads.box_predictor.bbox_pred.bias[FLOAT, 320]
    %560[FLOAT, 64x3x7x7]
    %561[FLOAT, 64]
    %563[FLOAT, 64x64x1x1]
    %564[FLOAT, 64]
    %566[FLOAT, 64x64x3x3]
    %567[FLOAT, 64]
    %569[FLOAT, 256x64x1x1]
    %570[FLOAT, 256]
    %572[FLOAT, 256x64x1x1]
    %573[FLOAT, 256]
    %575[FLOAT, 64x256x1x1]
    %576[FLOAT, 64]
    %578[FLOAT, 64x64x3x3]
    %579[FLOAT, 64]
    %581[FLOAT, 256x64x1x1]
    %582[FLOAT, 256]
    %584[FLOAT, 64x256x1x1]
    %585[FLOAT, 64]
    %587[FLOAT, 64x64x3x3]
    %588[FLOAT, 64]
    %590[FLOAT, 256x64x1x1]
    %591[FLOAT, 256]
    %593[FLOAT, 128x256x1x1]
    %594[FLOAT, 128]
    %596[FLOAT, 128x128x3x3]
    %597[FLOAT, 128]
    %599[FLOAT, 512x128x1x1]
    %600[FLOAT, 512]
    %602[FLOAT, 512x256x1x1]
    %603[FLOAT, 512]
    %605[FLOAT, 128x512x1x1]
    %606[FLOAT, 128]

```

%608[FLOAT, 128x128x3x3]
%609[FLOAT, 128]
%611[FLOAT, 512x128x1x1]
%612[FLOAT, 512]
%614[FLOAT, 128x512x1x1]
%615[FLOAT, 128]
%617[FLOAT, 128x128x3x3]
%618[FLOAT, 128]
%620[FLOAT, 512x128x1x1]
%621[FLOAT, 512]
%623[FLOAT, 128x512x1x1]
%624[FLOAT, 128]
%626[FLOAT, 128x128x3x3]
%627[FLOAT, 128]
%629[FLOAT, 512x128x1x1]
%630[FLOAT, 512]
%632[FLOAT, 256x512x1x1]
%633[FLOAT, 256]
%635[FLOAT, 256x256x3x3]
%636[FLOAT, 256]
%638[FLOAT, 1024x256x1x1]
%639[FLOAT, 1024]
%641[FLOAT, 1024x512x1x1]
%642[FLOAT, 1024]
%644[FLOAT, 256x1024x1x1]
%645[FLOAT, 256]
%647[FLOAT, 256x256x3x3]
%648[FLOAT, 256]
%650[FLOAT, 1024x256x1x1]
%651[FLOAT, 1024]
%653[FLOAT, 256x1024x1x1]
%654[FLOAT, 256]
%656[FLOAT, 256x256x3x3]
%657[FLOAT, 256]
%659[FLOAT, 1024x256x1x1]
%660[FLOAT, 1024]
%662[FLOAT, 256x1024x1x1]
%663[FLOAT, 256]
%665[FLOAT, 256x256x3x3]
%666[FLOAT, 256]
%668[FLOAT, 1024x256x1x1]
%669[FLOAT, 1024]
%671[FLOAT, 256x1024x1x1]
%672[FLOAT, 256]
%674[FLOAT, 256x256x3x3]
%675[FLOAT, 256]
%677[FLOAT, 1024x256x1x1]
%678[FLOAT, 1024]


```

%680[FLOAT, 256x1024x1x1]
%681[FLOAT, 256]
%683[FLOAT, 256x256x3x3]
%684[FLOAT, 256]
%686[FLOAT, 1024x256x1x1]
%687[FLOAT, 1024]
%689[FLOAT, 512x1024x1x1]
%690[FLOAT, 512]
%692[FLOAT, 512x512x3x3]
%693[FLOAT, 512]
%695[FLOAT, 2048x512x1x1]
%696[FLOAT, 2048]
%698[FLOAT, 2048x1024x1x1]
%699[FLOAT, 2048]
%701[FLOAT, 512x2048x1x1]
%702[FLOAT, 512]
%704[FLOAT, 512x512x3x3]
%705[FLOAT, 512]
%707[FLOAT, 2048x512x1x1]
%708[FLOAT, 2048]
%710[FLOAT, 512x2048x1x1]
%711[FLOAT, 512]
%713[FLOAT, 512x512x3x3]
%714[FLOAT, 512]
%716[FLOAT, 2048x512x1x1]
%717[FLOAT, 2048]
) {
%302 = AliasWithName[is_backward = 0, name = 'data'](%data)
%303 = AliasWithName[is_backward = 0, name = 'im_info'](%im_info)
%304 = Cast[to = 1](%302)
%305 = Constant[value = <Tensor>]()
%306 = Sub(%304, %305)
%307 = Constant[value = <Tensor>]()
%308 = Div(%306, %307)
%309 = AliasWithName[is_backward = 0, name = 'normalized_data'](%308)
%559 = Conv[dilations = [1, 1], group = 1, kernel_shape = [7, 7], pads = [3,
3, 3, 3], strides = [2, 2]](%309, %560, %561)
%312 = Relu(%559)
%313 = MaxPool[kernel_shape = [3, 3], pads = [1, 1, 1, 1], strides = [2,
2]](%312)
%562 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%313, %563, %564)
%316 = Relu(%562)
%565 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%316, %566, %567)
%319 = Relu(%565)
%568 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%319, %569, %570)

```

```

%571 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%313, %572, %573)
%324 = Add(%568, %571)
%325 = Relu(%324)
%574 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%325, %575, %576)
%328 = Relu(%574)
%577 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%328, %578, %579)
%331 = Relu(%577)
%580 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%331, %581, %582)
%334 = Add(%580, %325)
%335 = Relu(%334)
%583 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%335, %584, %585)
%338 = Relu(%583)
%586 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%338, %587, %588)
%341 = Relu(%586)
%589 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%341, %590, %591)
%344 = Add(%589, %335)
%345 = Relu(%344)
%592 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [2, 2]](%345, %593, %594)
%348 = Relu(%592)
%595 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%348, %596, %597)
%351 = Relu(%595)
%598 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%351, %599, %600)
%601 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [2, 2]](%345, %602, %603)
%356 = Add(%598, %601)
%357 = Relu(%356)
%604 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%357, %605, %606)
%360 = Relu(%604)
%607 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%360, %608, %609)
%363 = Relu(%607)
%610 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%363, %611, %612)
%366 = Add(%610, %357)
%367 = Relu(%366)
%613 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%367, %614, %615)

```

```

%370 = Relu(%613)
%616 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%370, %617, %618)
%373 = Relu(%616)
%619 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%373, %620, %621)
%376 = Add(%619, %367)
%377 = Relu(%376)
%622 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%377, %623, %624)
%380 = Relu(%622)
%625 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%380, %626, %627)
%383 = Relu(%625)
%628 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%383, %629, %630)
%386 = Add(%628, %377)
%387 = Relu(%386)
%631 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [2, 2]](%387, %632, %633)
%390 = Relu(%631)
%634 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%390, %635, %636)
%393 = Relu(%634)
%637 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%393, %638, %639)
%640 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [2, 2]](%387, %641, %642)
%398 = Add(%637, %640)
%399 = Relu(%398)
%643 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%399, %644, %645)
%402 = Relu(%643)
%646 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%402, %647, %648)
%405 = Relu(%646)
%649 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%405, %650, %651)
%408 = Add(%649, %399)
%409 = Relu(%408)
%652 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%409, %653, %654)
%412 = Relu(%652)
%655 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%412, %656, %657)
%415 = Relu(%655)
%658 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%415, %659, %660)

```

```

%418 = Add(%658, %409)
%419 = Relu(%418)
%661 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%419, %662, %663)
%422 = Relu(%661)
%664 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%422, %665, %666)
%425 = Relu(%664)
%667 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%425, %668, %669)
%428 = Add(%667, %419)
%429 = Relu(%428)
%670 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%429, %671, %672)
%432 = Relu(%670)
%673 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%432, %674, %675)
%435 = Relu(%673)
%676 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%435, %677, %678)
%438 = Add(%676, %429)
%439 = Relu(%438)
%679 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%439, %680, %681)
%442 = Relu(%679)
%682 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%442, %683, %684)
%445 = Relu(%682)
%685 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%445, %686, %687)
%448 = Add(%685, %439)
%449 = Relu(%448)
%688 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [2, 2]](%449, %689, %690)
%452 = Relu(%688)
%691 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%452, %692, %693)
%455 = Relu(%691)
%694 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%455, %695, %696)
%697 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [2, 2]](%449, %698, %699)
%460 = Add(%694, %697)
%461 = Relu(%460)
%700 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%461, %701, %702)
%464 = Relu(%700)
%703 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,

```

```

1, 1, 1], strides = [1, 1]](%464, %704, %705)
%467 = Relu(%703)
%706 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%467, %707, %708)
%470 = Add(%706, %461)
%471 = Relu(%470)
%709 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%471, %710, %711)
%474 = Relu(%709)
%712 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%474, %713, %714)
%477 = Relu(%712)
%715 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%477, %716, %717)
%480 = Add(%715, %471)
%481 = Relu(%480)
%482 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%481, %_wrapped_model.backbone.fpn_lateral5.weight,
%_wrapped_model.backbone.fpn_lateral5.bias)
%483 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%482, %_wrapped_model.backbone.fpn_output5.weight,
%_wrapped_model.backbone.fpn_output5.bias)
%484 = ResizeNearest[height_scale = 2, order = 'NCHW', width_scale = 2](%482)
%485 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%449, %_wrapped_model.backbone.fpn_lateral4.weight,
%_wrapped_model.backbone.fpn_lateral4.bias)
%486 = Add(%485, %484)
%487 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%486, %_wrapped_model.backbone.fpn_output4.weight,
%_wrapped_model.backbone.fpn_output4.bias)
%488 = ResizeNearest[height_scale = 2, order = 'NCHW', width_scale = 2](%486)
%489 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%387, %_wrapped_model.backbone.fpn_lateral3.weight,
%_wrapped_model.backbone.fpn_lateral3.bias)
%490 = Add(%489, %488)
%491 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%490, %_wrapped_model.backbone.fpn_output3.weight,
%_wrapped_model.backbone.fpn_output3.bias)
%492 = ResizeNearest[height_scale = 2, order = 'NCHW', width_scale = 2](%490)
%493 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%345, %_wrapped_model.backbone.fpn_lateral2.weight,
%_wrapped_model.backbone.fpn_lateral2.bias)
%494 = Add(%493, %492)
%495 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%494, %_wrapped_model.backbone.fpn_output2.weight,
%_wrapped_model.backbone.fpn_output2.bias)
%496 = MaxPool[kernel_shape = [1, 1], pads = [0, 0, 0, 0], strides = [2,
2]](%483)

```

```

    %497 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%495,
%_wrapped_model.proposal_generator.rpn_head.conv.weight,
%_wrapped_model.proposal_generator.rpn_head.conv.bias)
    %498 = Relu(%497)
    %499 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%498,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.weight,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.bias)
    %500 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%498,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.weight,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.bias)
    %501 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%491,
%_wrapped_model.proposal_generator.rpn_head.conv.weight,
%_wrapped_model.proposal_generator.rpn_head.conv.bias)
    %502 = Relu(%501)
    %503 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%502,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.weight,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.bias)
    %504 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%502,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.weight,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.bias)
    %505 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%487,
%_wrapped_model.proposal_generator.rpn_head.conv.weight,
%_wrapped_model.proposal_generator.rpn_head.conv.bias)
    %506 = Relu(%505)
    %507 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%506,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.weight,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.bias)
    %508 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%506,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.weight,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.bias)
    %509 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%483,
%_wrapped_model.proposal_generator.rpn_head.conv.weight,
%_wrapped_model.proposal_generator.rpn_head.conv.bias)
    %510 = Relu(%509)
    %511 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%510,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.weight,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.bias)

```

```

    %512 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%510,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.weight,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.bias)
    %513 = Conv[dilations = [1, 1], group = 1, kernel_shape = [3, 3], pads = [1,
1, 1, 1], strides = [1, 1]](%496,
%_wrapped_model.proposal_generator.rpn_head.conv.weight,
%_wrapped_model.proposal_generator.rpn_head.conv.bias)
    %514 = Relu(%513)
    %515 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%514,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.weight,
%_wrapped_model.proposal_generator.rpn_head.objectness_logits.bias)
    %516 = Conv[dilations = [1, 1], group = 1, kernel_shape = [1, 1], pads = [0,
0, 0, 0], strides = [1, 1]](%514,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.weight,
%_wrapped_model.proposal_generator.rpn_head.anchor_deltas.bias)
    %517, %518 = GenerateProposals[angle_bound_hi = 180, angle_bound_lo = -180,
angle_bound_on = 1, clip_angle_thresh = 1, legacy_plus_one = 0, min_size = 0,
nms_thresh = 0.699999988079071, post_nms_topN = 1000, pre_nms_topN = 1000,
spatial_scale = 0.25](%499, %500, %303,
%_wrapped_model.proposal_generator.anchor_generator.cell_anchors.0)
    %519, %520 = GenerateProposals[angle_bound_hi = 180, angle_bound_lo = -180,
angle_bound_on = 1, clip_angle_thresh = 1, legacy_plus_one = 0, min_size = 0,
nms_thresh = 0.699999988079071, post_nms_topN = 1000, pre_nms_topN = 1000,
spatial_scale = 0.125](%503, %504, %303,
%_wrapped_model.proposal_generator.anchor_generator.cell_anchors.1)
    %521, %522 = GenerateProposals[angle_bound_hi = 180, angle_bound_lo = -180,
angle_bound_on = 1, clip_angle_thresh = 1, legacy_plus_one = 0, min_size = 0,
nms_thresh = 0.699999988079071, post_nms_topN = 1000, pre_nms_topN = 1000,
spatial_scale = 0.0625](%507, %508, %303,
%_wrapped_model.proposal_generator.anchor_generator.cell_anchors.2)
    %523, %524 = GenerateProposals[angle_bound_hi = 180, angle_bound_lo = -180,
angle_bound_on = 1, clip_angle_thresh = 1, legacy_plus_one = 0, min_size = 0,
nms_thresh = 0.699999988079071, post_nms_topN = 1000, pre_nms_topN = 1000,
spatial_scale = 0.03125](%511, %512, %303,
%_wrapped_model.proposal_generator.anchor_generator.cell_anchors.3)
    %525, %526 = GenerateProposals[angle_bound_hi = 180, angle_bound_lo = -180,
angle_bound_on = 1, clip_angle_thresh = 1, legacy_plus_one = 0, min_size = 0,
nms_thresh = 0.699999988079071, post_nms_topN = 1000, pre_nms_topN = 1000,
spatial_scale = 0.015625](%515, %516, %303,
%_wrapped_model.proposal_generator.anchor_generator.cell_anchors.4)
    %rpn_rois = CollectRpnProposals[rpn_max_level = 6, rpn_min_level = 2,
rpn_post_nms_topN = 1000](%517, %519, %521, %523, %525, %518, %520, %522, %524,
%526)
    %528, %529, %530, %531, %532 = DistributeFpnProposals[legacy_plus_one = 0,
roi_canonical_level = 4, roi_canonical_scale = 224, roi_max_level = 5,
roi_min_level = 2](%rpn_rois)

```

```

    %533 = RoIAlign[aligned = 1, order = 'NCHW', pooled_h = 7, pooled_w = 7,
sampling_ratio = 0, spatial_scale = 0.25](%495, %528)
    %534 = RoIAlign[aligned = 1, order = 'NCHW', pooled_h = 7, pooled_w = 7,
sampling_ratio = 0, spatial_scale = 0.125](%491, %529)
    %535 = RoIAlign[aligned = 1, order = 'NCHW', pooled_h = 7, pooled_w = 7,
sampling_ratio = 0, spatial_scale = 0.0625](%487, %530)
    %roi_feat_fpn = RoIAlign[aligned = 1, order = 'NCHW', pooled_h = 7, pooled_w =
7, sampling_ratio = 0, spatial_scale = 0.03125](%483, %531)
    %537 = Concat[axis = 0](%533, %534, %535, %roi_feat_fpn)
    %538 = BatchPermutation(%537, %532)
    %539 = Flatten[axis = 1](%538)
    %540 = Gemm[alpha = 1, beta = 1, transB = 1](%539,
%_wrapped_model.roi_heads.box_head.fc1.weight,
%_wrapped_model.roi_heads.box_head.fc1.bias)
    %541 = Relu(%540)
    %542 = Gemm[alpha = 1, beta = 1, transB = 1](%541,
%_wrapped_model.roi_heads.box_head.fc2.weight,
%_wrapped_model.roi_heads.box_head.fc2.bias)
    %543 = Relu(%542)
    %544 = Gemm[alpha = 1, beta = 1, transB = 1](%543,
%_wrapped_model.roi_heads.box_predictor.cls_score.weight,
%_wrapped_model.roi_heads.box_predictor.cls_score.bias)
    %545 = Gemm[alpha = 1, beta = 1, transB = 1](%543,
%_wrapped_model.roi_heads.box_predictor.bbox_pred.weight,
%_wrapped_model.roi_heads.box_predictor.bbox_pred.bias)
    %546 = Softmax[axis = 1](%544)
    %547 = Concat[axis = 0](%rpn_rois)
    %548, %549 = BBoxTransform[angle_bound_hi = 180, angle_bound_lo = -180,
angle_bound_on = 1, apply_scale = 1, clip_angle_thresh = 1, legacy_plus_one = 0,
rotated = 0, weights = [10, 10, 5, 5]](%547, %545, %303)
    %roi_score_nms, %roi_bbox_nms.1, %roi_class_nms, %roi_batch_splits_nms,
%roi_keeps_nms, %roi_keeps_size_nms = BoxWithNMSLimit[cls_agnostic_bbox_reg = 0,
detections_per_im = 100, input_boxes_include_bg_cls = 0, legacy_plus_one = 0,
nms = 0.5, output_classes_include_bg_cls = 0, rotated = 0, score_thresh = 0.5,
soft_nms_enabled = 0, soft_nms_method = 'linear', soft_nms_min_score_thres =
0.001000000004749745, soft_nms_sigma = 0.5](%546, %548, %549)
    %556 = AliasWithName[is_backward = 0, name = 'class_nms'](%roi_class_nms)
    %557 = AliasWithName[is_backward = 0, name = 'score_nms'](%roi_score_nms)
    %roi_bbox_nms = AliasWithName[is_backward = 0, name =
'bbox_nms'](%roi_bbox_nms.1)
    return %roi_bbox_nms, %557, %556
}

```

```

[12]: from detectron2.export.caffe2_modeling import
      ↪ convert_batched_inputs_to_c2_format, META_ARCH_CAFFE2_EXPORT_TYPE_MAP

      # image preparation from `DefaultPredictor`

```



```
# https://github.com/facebookresearch/detectron2/blob/master/detectron2/engine/defaults.py#L214
original_image = cv2.imread('datasets/coco/val2017/000000023937.jpg')

aug = T.ResizeShortestEdge(
    [cfg.INPUT.MIN_SIZE_TEST, cfg.INPUT.MIN_SIZE_TEST], cfg.INPUT.MAX_SIZE_TEST
)

height, width = original_image.shape[:2]
image = aug.get_transform(original_image).apply_image(original_image)
image = torch.as_tensor(image.astype("float32").transpose(2, 0, 1))

inputs = {"image": image, "height": height, "width": width}
(image, img_info) = convert_batched_inputs_to_c2_format([inputs], 32, "cpu")
```

```
[13]: print(model.graph.input[0])
      print(model.graph.input[1])

# two inputs are needed `data` wich is the image and `img_info` wich represents
↳ the shape of the image
W = {
    model.graph.input[0].name: image.data.numpy(),
    model.graph.input[1].name: img_info.data.numpy()
}
```

```
name: "data"
type {
  tensor_type {
    elem_type: 2
    shape {
      dim {
        dim_value: 1
      }
      dim {
        dim_value: 3
      }
      dim {
        dim_value: 800
      }
      dim {
        dim_value: 1216
      }
    }
  }
}
```

```
name: "im_info"
```

```

type {
  tensor_type {
    elem_type: 1
    shape {
      dim {
        dim_value: 1
      }
      dim {
        dim_value: 3
      }
    }
  }
}

```

```

[14]: # run onnx inference
rep = backend.prepare(model)
raw_onnx_outputs = rep.run(W)

```

```

/home/users/fassbinb/.local/lib/python3.8/site-
packages/caffe2/python/onnx/backend.py:690: UserWarning: Unrecognized operator
set org.pytorch._caffe2
  warnings.warn("Unrecognized operator set {}".format(imp.domain))

```

```

[15]: onnx_outputs = {
      "bbox_nms": torch.tensor(raw_onnx_outputs.roi_bbox_nms),
      "score_nms": torch.tensor(raw_onnx_outputs._1),
      "class_nms": torch.tensor(raw_onnx_outputs._2)
    }

```

```

[16]: # convert output to detectron2 compatible output
from detectron2.modeling import meta_arch
from detectron2.export.caffe2_modeling import assemble_rcnn_outputs_by_name

image_sizes = [[int(im[0]), int(im[1])] for im in img_info]
results = assemble_rcnn_outputs_by_name(image_sizes, onnx_outputs)

# replace with the used model
outputs = getattr(meta_arch, cfg.MODEL.META_ARCHITECTURE)._postprocess(results,
↳ [inputs], image_sizes)

```

```

[17]: v = Visualizer(original_image[:, :, ::-1], MetadataCatalog.get(cfg.DATASETS.
↳ TRAIN[0]), scale=1.2)
out = v.draw_instance_predictions(outputs[0]["instances"].to("cpu"))
cv2_imshow(out.get_image()[:, :, ::-1])

```

